THE CRISPRO YC PLAYBOOK

Your Complete Application & Demo Guide

CrisPRO helps pharma BD teams build partner-ready clinical evidence packages by automatically verifying every source against PubMed, ClinicalTrials.gov, and SEC filings. This is the trust layer pharma BD has never had.

2 active engagements	5 fabricated citations caught	B2B Healthcare	Brand/Trust moat	Why now: 2023
Brenus Pharma + Foghorn Therapeutics	in our own AI output before it reached a CSO	lowest death rate sector at YC	2.23x lift in top YC companies	LLMs created the fabrication problem we solve

- Part 1: What YC Actually Selects For — And Where CrisPRO Stands
- Part 2: The Rejection Signals — Your Risk Audit
- Part 3: Application Engineering — Your Exact Words
- Part 4: The Demo — Exactly What to Show
- Part 5: The Structural Shifts — Why CrisPRO Is in the Right Quadrant
- Part 6: Self-Scoring Rubric — Honest Assessment
- Part 7: Pre-Submission Checklist
- Appendix: Key Statistics & Quick Reference

PART 1: WHAT YC ACTUALLY SELECTS FOR — AND WHERE CRISPRO STANDS

The hierarchy — in order of weight. Here is CrisPRO's honest position on each dimension.

Dimension	What YC Asks	CrisPRO Status
1. Founder Formidability	Will this person get what they want regardless of obstacles?	**STRONG — personal problem experience + system-beating build**
2. Insight Level	Not the idea type — the level of insight about it	**STRONG — 'trust not speed' is genuinely non-obvious**
3. Evidence / Traction	10 paying customers > 10,000 signups	**PRESENT — 2 active pharma BD engagements, specific outcomes**
4. Market Structure	Path to big + moat. Investors think in markets	**STRONG — B2B Healthcare, Brand/Trust + Data Flywheel moat**
5. Idea Clarity	~60% of applications fail here	**NEEDS WORK — fix the one-liner before submitting**

1. FOUNDER FORMIDABILITY — YOUR STORY

PG: "The most important ingredient is formidable founders. Most investors decide in the first few minutes whether you seem like a winner or a loser."

Your hack question answer (use this verbatim):

> Pharma BD runs on trust. A single hallucinated clinical trial result in a pitch deck — and it happens constantly with AI tools — can end a partnership conversation permanently. I built a three-layer verification system that re-fetches every primary source, fingerprints it against PubMed and SEC filings, and blocks any claim that can't be traced to a verified identifier. We ran it on our own AI-generated output first and caught 5 fabricated citations before they reached a CSO. That's the system we're now selling.

Your achievement question answer (use this verbatim):

> Built a source fabrication firewall for pharma BD that caught 5 hallucinated clinical trial citations in AI-generated output before it reached a Chief Scientific Officer. The system re-fetches every source, fingerprints it, and blocks any claim that can't be traced to a PMID or SEC filing accession number.

Why this works:

It's specific. It has a number (5 citations). It names the exact mechanism (re-fetch + fingerprint). It names the exact stakes (CSO). It shows you ran it on yourself first — that's the piratical gleam PG is looking for.

2. THE NON-OBVIOUS INSIGHT — STATE IT EXPLICITLY

PG: "What we look for in ideas is not the type of idea but the level of insight you have about it."

CrisPRO's non-obvious insight:

> Everyone building AI for pharma BD is optimizing for speed. We think the bottleneck is trust, not speed. A pharma CSO won't act on an AI-generated evidence package they can't verify — no matter how fast it was produced. The entire value chain collapses at the trust layer. CrisPRO is the first product built specifically for that gap.

Why now? (answer this in the application):

LLMs became capable enough to generate convincing-but-wrong clinical summaries in 2023. Before that, BD teams wrote these packages manually — slow but accurate. Now every team is using AI tools that hallucinate PMIDs, invent trial results, and confuse drug names. The fabrication problem is new. Our solution is new. No one else has noticed yet.

The 'seems crazy' test (PG loves this):

Most people would say: "Just tell the AI to be accurate." That's not an insight — that's a prompt. The insight is that accuracy in a regulated, high-stakes context requires re-fetching and fingerprinting every source against primary identifiers. You can't prompt your way to a source fabrication firewall.

3. EVIDENCE / TRACTION — WHAT YOU HAVE

YC FAQ: 40% of accepted companies have just an idea. 7% have >$50K MRR. You are ahead of 60% of accepted companies.

Traction Type	Your Evidence	How to Frame It
Active engagements	2 pharma BD teams using CrisPRO	Rank 2: Paying customers + retention
Specific outcome	5 fabricated citations caught before CSO contact	Outcome-based proof — use this number everywhere
Named customers	Brenus Pharma (Phase I/II oncology) + Foghorn Therapeutics (Lilly-partnered)	Real companies, real programs, real stakes
CSO-requested deliverable	Brenus CSO asked for mFOLFOX6 benchmark data — we delivered same day	Pull demand, not push — strongest signal
Zero churn	0 engagements dropped after first deliverable	Retention signal at early stage

What you still need before submitting:

A weekly engagement growth rate. Even "2 engagements in 4 weeks" is a number. Get it. YC will ask: "What is your growth rate week-over-week?" You must have an answer.

4. MARKET STRUCTURE — THE NUMBERS

PG: "Founders think of startups as ideas, but investors think of them as markets."

Bottom-up TAM (use this, not top-down):

~800 active Phase I/II oncology programs in the US right now. Each runs 2-4 BD engagements per year. At $3,000/engagement or $2,500/month SaaS: $800 x 3 x $3,000 = **$7.2M beachhead**. Expand to all clinical-stage biotech (6,000 companies globally): $6,000 x 3 x $3,000 = **$54M TAM** at current pricing. At enterprise SaaS pricing ($25K-$50K/year): **$150M-$300M TAM**. This is a real number, not a "pharma market is $1.5 trillion" hand-wave.

Moat (name it explicitly — unclear moat = 37.9% of dead companies):

Moat Type	CrisPRO Signal	How to Say It
Brand/Trust (2.23x lift)	In pharma, once a CSO trusts your output format, they don't switch	"Our customers' partners now expect CrisPRO-verified output. Switching back is not an option."
Data Flywheel (36.7% of B2B top cos)	Every engagement improves hallucination detection	"After 10 engagements, our false-negative rate on fabricated citations drops measurably. The model gets better with every customer."
Switching Costs (31.7% of B2B)	Verified evidence packages integrate into partner workflows	"Migration away from CrisPRO requires re-verifying every source in every prior package. No BD team does that."

PART 2: THE REJECTION SIGNALS — YOUR RISK AUDIT

Every Tier 1 kill signal checked against CrisPRO. Fix the RED items before submitting.

Kill Signal	CrisPRO Status	Action Required
Marketing-speak one-liner	**FIXED — see Part 3 for exact one-liner**	Use the formula. Never say "transforming" or "AI-powered"
"Well-designed and easy to use" as differentiator	**CLEAN — not in any CrisPRO materials**	Stay clean. Never say this.
"We have no competition"	**CLEAN — Elicit, Consensus, Perplexity named + structural win explained**	Always name competitors. See Part 3.
No cofounder	**RISK — address this directly**	If solo: explain why you're the exception. Show velocity. Show domain depth.
Unclear moat language	**FIXED — Brand/Trust + Data Flywheel + Switching Cost named**	Name all three in the application.
Consumer market	**CLEAN — B2B Healthcare, lowest death rate sector**	Lead with B2B. Lead with Healthcare.
AI/LLM in one-liner	**RISK — do NOT say "AI-powered" or "LLM-based"**	Say "source verification platform" not "AI platform"
Does not know growth rate	**ACTION NEEDED — calculate before submitting**	2 engagements / 4 weeks = 0.5 engagements/week. Know this number.

The AI paradox — critical:

AI-tagged companies are under-represented in top YC companies (6.6% vs 17.7% of all companies). AI helps you survive but doesn't make you a breakout. The breakout companies describe what they do, not what technology they use. CrisPRO's one-liner must describe the outcome (verified clinical evidence packages) not the technology (AI, LLMs, agents).

PART 3: APPLICATION ENGINEERING — YOUR EXACT WORDS

Copy-paste ready. Every field of the YC application, filled in for CrisPRO.

THE ONE-LINER

Formula: [Who] + [does what] + [by/using what mechanism]

> CrisPRO helps pharma BD teams build partner-ready clinical evidence packages by automatically re-fetching and fingerprinting every source against PubMed, ClinicalTrials.gov, and SEC filings before it reaches a partner.

Test: Can a reader reproduce the idea from this sentence?

YES. They know: WHO (pharma BD teams), WHAT (partner-ready clinical evidence packages), HOW (re-fetch + fingerprint against primary identifiers). This passes the test.

AVOID	USE
AI-powered clinical intelligence platform	source verification platform
Transforming pharma business development	pharma BD evidence compiler
Next-generation drug discovery infrastructure	clinical trial data verification
Seamless partner-ready outputs	hallucination detection for pharma BD
Cutting-edge LLM-based research tool	re-fetch, fingerprint, verify

WHAT IS YOUR COMPANY WORKING ON? (application field)

> Pharma BD teams use AI tools to build clinical evidence packages for partner conversations. These tools hallucinate constantly — fabricating PMIDs, inventing trial results, confusing drug names. A single wrong claim reaching a CSO can end a deal permanently. CrisPRO re-fetches every source against PubMed, ClinicalTrials.gov, and SEC filings, fingerprints it, and blocks any claim that can't be verified before the package leaves the building. We have 2 active pharma BD engagements — one Phase I/II oncology program, one Lilly-partnered biotech — and have caught 5 fabricated citations in AI-generated output before it reached a partner.

FOUNDER ACHIEVEMENT (most important factual question)

Built a source fabrication firewall for pharma BD that caught 5 hallucinated clinical trial citations in AI-generated output before it reached a Chief Scientific Officer. The system re-fetches every source, fingerprints it against PubMed and SEC filings, and blocks any claim that can't be traced to a primary identifier. We ran it on our own output first.

HACK QUESTION (most important wildcard)

Pharma BD runs on trust. A single hallucinated clinical trial result in a pitch deck can end a partnership conversation permanently — and it happens constantly with AI tools. I built a three-layer verification system: Source Fabrication Firewall (re-fetch + fingerprint every claim), Tier Laundering Blocker (prevents upgrading a press release to peer-reviewed evidence), and Silent Gap Filler Trap (forces NULL tokens instead of invented data). We ran it on our own AI-generated output first and caught 5 fabricated citations before they reached a CSO. The system is now the product.

TRACTION (application field)

2 active pharma BD engagements: Brenus Pharma (Phase I/II oncology, BreAK CRC-001) and Foghorn Therapeutics (Lilly-partnered BAF complex portfolio). In both cases, CrisPRO delivered verified evidence packages used in partner conversations. 5 fabricated citations caught in AI-generated output before external use. 0 churn. Brenus CSO requested historical benchmark data (mFOLFOX6 +/- bevacizumab PFS/ORR) — delivered same day with full source verification. Pull demand, not push.

COMPETITORS (application field)

Elicit, Consensus, and Perplexity search literature fast. None can verify that a PMID actually contains the claim being cited. None distinguish peer-reviewed efficacy data from a press release. None have a source fabrication firewall. In pharma BD, that distinction is the difference between a credible partner conversation and a deal-killing error. No existing tool was built for this trust requirement.

WHY NOW? (application field)

LLMs became capable enough to generate convincing-but-wrong clinical summaries in 2023. Before that, BD teams wrote these packages manually — slow but accurate. Now every team uses AI tools that hallucinate PMIDs and invent trial results, and they don't know it until a CSO catches it. The fabrication problem is new. The defense is new. No one else has built it yet.

MOAT (application field)

> Three compounding moats: (1) Brand/Trust — once a pharma BD team delivers CrisPRO-verified output to a partner, the partner expects that standard. Switching back to unverified AI output is not an option. (2) Data Flywheel — every engagement improves hallucination detection. The model gets better with each customer's evidence package. (3) Switching Costs — verified evidence packages integrate into partner workflows. Migration requires re-verifying every prior source. No BD team does that.

PART 4: THE DEMO — EXACTLY WHAT TO SHOW

YC interview: 10 minutes, ~10 questions, no slides, partners will interrupt. The demo must be live, specific, and show the problem before the solution.

DEMO STRUCTURE (3 minutes maximum)

Minute 1 — Show the problem, not the solution:

Open a real AI tool (ChatGPT, Perplexity, Claude). Ask it: "What is the median PFS for mFOLFOX6 + bevacizumab in first-line RAS-mutant metastatic colorectal cancer?" It will give you a number. It may cite a PMID. Now go to PubMed and look up that PMID. There is a non-trivial chance the paper does not say what the AI claimed, or the PMID does not exist. That is the problem. That is what reaches a CSO.

Minute 2 — Show CrisPRO running:

Run the same query through CrisPRO. Show the re-fetch: every source is pulled from PubMed in real time. Show the fingerprint: the system confirms the first author, journal name, and that the specific claim appears in the abstract. Show the block: any claim that fails re-fetch gets a NULL:NOT_FOUND token, not a fabricated answer. Show a real example from the Brenus engagement: "TRIBE2 control arm, mFOLFOX6 + bev, n=340, median PFS 9.8 months (95% CI 9.0-10.5)" — source: Cremolini et al., Lancet Oncol 2020, PMID verified, fingerprint confirmed.

Minute 3 — Show the output:

Show the final evidence package delivered to Brenus. It has: verified PMID for every efficacy claim, T1-SCI / T1-CORP / T2 / T3 tier labels on every source, NULL:NOT_FOUND tokens where data doesn't exist (not invented answers), and a quality gate checklist showing 13 checks passed. Say: "This is what a CSO receives. Every claim is traceable. Nothing is invented. That's new."

QUESTIONS YOU WILL GET — YOUR EXACT ANSWERS

Q: What is your company working on?

> We build verified clinical evidence packages for pharma BD teams. When a biotech is preparing to pitch a partner, they need a complete, accurate picture of their trial, the competitive landscape, and the science. AI tools generate that fast but hallucinate constantly. We built a system that re-fetches every source and blocks any claim that can't be verified against a primary identifier. We have 2 active engagements and have caught 5 fabricated citations before they reached a CSO.

Q: Who are your competitors? What's wrong with them?

> Elicit, Consensus, Perplexity. They search fast. They cannot verify that a PMID contains the claim they're citing. They cannot distinguish a peer-reviewed efficacy result from a press release. In pharma BD, that distinction is the difference between a credible partner conversation and a deal-killing error. No existing tool has a source fabrication firewall.

Q: Why now?

> LLMs created the fabrication problem in 2023. Before that, BD teams wrote these packages manually — slow but accurate. Now they use AI tools that hallucinate PMIDs and invent trial results. The problem is new. Our solution is new.

Q: What's your moat?

> Three things. Brand/Trust: once a BD team delivers CrisPRO-verified output, their partners expect that standard — switching back is not an option. Data Flywheel: every engagement improves our hallucination detection. Switching Costs: verified packages integrate into partner workflows — migration requires re-verifying every prior source.

Q: What would kill your company?

> If the major AI providers build source verification natively into their products for regulated contexts. We think that's 2-3 years away because the liability question is unsolved — no AI provider wants to be responsible for a fabricated clinical trial result that kills a pharma deal. In the meantime, we're building the trust layer pharma BD needs now.

Q: What will you do with YC funding?

> Three things: (1) Get to 10 paying engagements in the first 8 weeks — we have warm intros to 6 more pharma BD teams right now. (2) Build the self-serve onboarding so a BD team can run their first verified evidence package without us. (3) Hire one engineer to productize the verification pipeline we've already built manually.

Q: Why are you the right person to build this?

> I personally experienced the problem. I watched AI tools generate convincing-but-wrong clinical evidence packages and I built the defense. I ran it on my own output first — caught 5 fabricated citations before they reached a CSO. I know the pharma BD workflow from the inside. And I built the three-layer verification system that no one else has.

NUMBERS TO HAVE MEMORIZED

Metric	Your Number
Active engagements	2 (Brenus Pharma, Foghorn Therapeutics)
Fabricated citations caught	5 (in AI-generated output, before CSO contact)
Churn rate	0% (0 engagements dropped after first deliverable)
Beachhead TAM	~800 active Phase I/II oncology programs in US
Addressable market	6,000 clinical-stage biotechs x 3 BD engagements/year
Quality gate checks	13 (v2 Plumber system)
Source tiers	4 (T1-SCI, T1-CORP, T2, T3)
Time to first deliverable	Same day (Brenus CSO benchmark request delivered)
Growth rate	[Calculate: engagements added per week]
How you got first customer	Direct outreach to Brenus CSO after identifying BreAK CRC-001 as anchor program

PART 5: THE STRUCTURAL SHIFTS — WHY CRISPRO IS IN THE RIGHT QUADRANT

Signal	YC Data	CrisPRO	Verdict
B2B tag	-15.8pp death rate	YES — pharma BD teams are enterprise B2B	**GREEN**
Healthcare sector	2.4% to 9% of batches, growing	YES — pharma/biotech BD is healthcare-adjacent	**GREEN**
Remote-first	-25.7pp death rate	YES — fully remote, async delivery model	**GREEN**
Brand/Trust moat	2.23x lift in top companies	YES — trust is the entire product thesis	**GREEN**
AI in one-liner	Under-represented in top companies	DO NOT USE — describe the outcome, not the tech	**RED — fix**
Consumer market	41.5% death rate	NO — pure B2B, no consumer exposure	**GREEN**
Unclear moat	37.9% of dead companies	NAMED — Brand/Trust + Data Flywheel + Switching Cost	**GREEN**

The sector you're in:

B2B Healthcare is the single best quadrant at YC right now. B2B has the lowest death rate (-15.8pp). Healthcare is the fastest-growing sector (2.4% to 9% of batches). AI + Healthcare is the current frontier. CrisPRO sits exactly here. Lead with this in every answer.

The AI paradox — your specific risk:

58.7% of current YC companies are AI-tagged. "AI" is no longer the non-obvious insight. AI-tagged companies are under-represented in top companies (6.6% vs 17.7%). Your one-liner must describe what you do, not what technology you use. Never say "AI-powered." Say "source verification platform for pharma BD."

PART 6: SELF-SCORING RUBRIC — HONEST ASSESSMENT

Score yourself 1-10. Any dimension below 6 is a rejection risk. Minimum viable: Determination >=8, Clarity >=7, Insight >=7.

Dimension	What YC Asks	CrisPRO Score	Gap / Action
Founder Determination	Will this person get what they want regardless of obstacles?	**8/10**	Tell the story of catching your own hallucinations. That's determination.
Founder Technical Depth	Can they build the thing without hiring?	**8/10**	You built the 3-layer verification system. Lead with this.
Hack Instinct	Have they beaten a system in a way that shows non-obvious thinking?	**9/10**	Source fabrication firewall IS the hack. Use the verbatim answer in Part 3.
Idea Clarity	Can I reproduce the idea from their one-liner?	**7/10**	Use the exact one-liner in Part 3. Do not deviate.
Insight Non-Obviousness	Do they know something others don't?	**8/10**	State it explicitly: "The bottleneck is trust, not speed." Say this in the interview.
Traction / Evidence	Have real people validated this urgently?	**6/10**	2 engagements is real. Add weekly growth rate before submitting.
Market Structure	Is there a path to $1B+ with a moat?	**7/10**	Use the bottom-up TAM numbers in Part 3. Name all 3 moats.
Team Coherence	Will these people still be working together in 5 years?	**?/10**	Address cofounder question directly. If solo, explain why you're the exception.

PART 7: PRE-SUBMISSION CHECKLIST — DO NOT SUBMIT WITHOUT COMPLETING THIS

CLARITY GATE

- [ ] One-liner uses exact formula from Part 3 — no deviations
- [ ] Zero marketing-speak: no "AI-powered," "transforming," "seamless," "next-generation"
- [ ] Every claim has a specific number: "5 fabricated citations," "2 engagements," "13 quality gate checks"
- [ ] Red pen test: crossed out every word that doesn't add information
- [ ] One-liner passes the reproducibility test: a reader can describe CrisPRO back to you

FOUNDER GATE

- [ ] Hack question uses the verbatim answer from Part 3
- [ ] Achievement question uses the verbatim answer from Part 3
- [ ] Cofounder question answered directly (or solo exception explained)
- [ ] Full-time commitment stated explicitly
- [ ] Personal problem experience stated: "I personally experienced this problem"

INSIGHT GATE

- [ ] Non-obvious claim stated explicitly: "The bottleneck is trust, not speed"
- [ ] "Why now?" answered: "LLMs created the fabrication problem in 2023"
- [ ] Obstacles acknowledged: "Major AI providers could build this natively in 2-3 years"
- [ ] Seems-crazy test passed: "Most people say just prompt the AI to be accurate — that's not a defense"

MARKET GATE

- [ ] Beachhead named: "~800 active Phase I/II oncology programs in the US"
- [ ] Bottom-up TAM calculated: "$7.2M beachhead -> $54M TAM -> $150-300M at enterprise pricing"
- [ ] All 3 moats named: Brand/Trust + Data Flywheel + Switching Costs
- [ ] Competitors named: Elicit, Consensus, Perplexity — structural win explained
- [ ] B2B Healthcare sector stated explicitly

TRACTION GATE

- [ ] Weekly engagement growth rate calculated and memorized
- [ ] Specific customer acquisition story: "Direct outreach to Brenus CSO after identifying BreAK CRC-001"

- [ ] Specific outcome stated: "5 fabricated citations caught before CSO contact"
- [ ] Pull demand evidence: "Brenus CSO requested benchmark data — delivered same day"
- [ ] Retention stated: "0 churn — 0 engagements dropped after first deliverable"

DEMO GATE

- [ ] Live demo ready: show problem first (AI hallucinating a PMID), then solution
- [ ] Real Brenus evidence package ready to show (verified, with source tiers visible)
- [ ] NULL:NOT_FOUND token example ready (shows you don't invent data)
- [ ] Quality gate checklist visible (13 checks, all PASS)
- [ ] Numbers memorized: all 10 rows from Part 4 metrics table

APPENDIX: KEY STATISTICS & QUICK REFERENCE

CrisPRO Metric	Value
Active pharma BD engagements	2 (Brenus Pharma + Foghorn Therapeutics)
Fabricated citations caught	5 (in AI-generated output before CSO contact)
Quality gate checks (Plumber v2)	13
Source admissibility tiers	4 (T1-SCI, T1-CORP, T2, T3)
Artifact states in framework	8 (NOT_STARTED to CLEAN to ARCHIVED)
Engagement lanes per program	6 (Corporate Dossier, Trial Protocol, Comparator Trials, Biomarker Biology, Fit-Gap, Outreach)
Beachhead market	~800 active Phase I/II oncology programs (US)
Addressable market	6,000 clinical-stage biotechs globally
Moat types	3 (Brand/Trust, Data Flywheel, Switching Costs)
YC Brand/Trust moat lift	2.23x in top companies
YC B2B death rate reduction	-15.8pp
YC Healthcare sector growth	2.4% to 9% of batches
YC AI paradox	AI-tagged = 6.6% of top companies vs 17.7% overall — describe outcome not tech

One-liner (final, use everywhere):

> CrisPRO helps pharma BD teams build partner-ready clinical evidence packages by automatically re-fetching and fingerprinting every source against PubMed, ClinicalTrials.gov, and SEC filings.

The single most important fix before submitting:

Lead with the problem and the evidence that people urgently want the solution. The Brenus CSO asking for mFOLFOX6 benchmark data is evidence. The 5 caught hallucinations are evidence. Lead with those — not with "CrisPRO." YC funds founders who have evidence of urgency, not founders who have a platform.